

Beyond Turing: Design, Attack, and Formalize a Test of Intelligence


Assessing Intelligence through Dynamic Adaptation and Deduction

Francesco Tartaglia

The Limits of Imitation: Communicating is not Comprehending

The Limitations of the Turing Test: Why Imitating Does Not Mean Reasoning

- The Imitation Game evaluates intelligence through linguistic deception in a purely textual and abstract environment, reflecting a strong anthropocentric bias (it measures how "human" the machine appears).
- This approach causes the test to reward the ability to "fake" human behavior, completely neglecting actual deductive and logical reasoning capabilities.
- Today we know that a system can produce linguistically perfect responses (like Stochastic Parrots) based solely on statistical models, without having any real cognition of what it is saying.



Conclusion: "A system that communicates excellently is not necessarily a system that thinks."

The Proposed Alternative: Zero-Shot Deduction Test (ZSDT)

Evaluating intelligence through adaptation and deduction in unknown environments

- Unlike the Turing Test, the ZSDT evaluates an agent placed in a totally unknown virtual environment, without being able to rely on any prior knowledge (**Zero-Shot approach**).
- The rules of the world are never explained to the agent in advance and are procedurally generated (they change with each new execution of the test).
- The agent must actively interact with the environment and map the consequences of its actions by observing feedback (e.g., "valid move", "invalid move", "score variation").
- The actual goal is to use this information to deduce the hidden logic and reach the objective (win the game) within a limited number of attempts.
- It analyzes the ability to operate in a partially observable environment, demonstrating true adaptability and inductive reasoning (learning by doing), overcoming the limits of simple data memorization.



What Does the ZSDT Actually Measure?

Beyond imitation: evaluating pure deductive reasoning and adaptability

- **Pure Inductive Reasoning:** Unlike Turing, the ZSDT does not evaluate whether the machine "appears human". It measures its ability to extrapolate general logical rules starting from a few direct interactions (trial and error).
- **Dynamic Problem-Solving:** The test evaluates how the agent uses the instant feedback of the environment to correct its own strategies in real-time, demonstrating true cognitive flexibility.
- **Elimination of "Training Data Bias":** Since the rules are procedurally generated and never seen before (Zero-Shot), the test prevents the machine from cheating using pre-memorized databases. It eliminates the "Stochastic Parrot" effect.
- **Objective (Non-Subjective) Intelligence:** The Turing Test relies on the opinions of a human judge (easily deceived). The ZSDT relies on mathematical and objective metrics: number of failed attempts, adaptation time, and achievement of the goal.

Hacking the ZSDT: Stress Testing and Failure Models

Identifying the limits: the two failure models of the test

- **The Need for Stress Testing:** Every rigorous evaluation test needs to be pushed to the limit to define its own vulnerabilities. By subjecting the ZSDT to a "stress test", we have identified two failure models that alter the validity of the intelligence measurement.
- **Failure Model 1 – The False Positive** (Passes the test, but is not intelligent): A "Brute Force" algorithm or a very fast random search bot. It could guess the sequence of winning moves by pure chance or by calculating millions of combinations per second, passing the test without having "understood" the logic in the slightest.
- **Failure Model 2 – The False Negative** (Fails the test, but is intelligent): A highly advanced large language model (LLM) or a pure mathematical AI. Despite possessing formidable theoretical intelligence, it would fail the test miserably because its architecture is not designed to interact with a spatial environment and deduce "gameplay" rules.

Formalizing the Agent: The PEAS Model for ZSDT

Concrete examples and descriptions of the architecture and environment

Agent

A reasoning software (an autonomous agent) integrated into the simulation environment, lacking pre-programmed instructions on the rules of the game.

Environment

- **Virtual Grid Map:** A digital maze or chessboard whose physical laws and traversal rules are procedurally generated.
- **Characteristics:** Partially observable (the agent sees the map but not the underlying logic), dynamic (rules can evolve), discrete, and deterministic.

Sensors

- **Spatial Sensor:** Receives the current position of the agent and the visual state of adjacent tiles (X, Y coordinates).
- **Feedback Sensor:** Receives the immediate result of each action from the system (Valid Move, Obstacle, Score Variation, Game Over).

Actuators

- **Discrete Movement Commands:** Invocation of actions such as NORTH, SOUTH, EAST, WEST.
- **Interaction Commands:** State modification actions such as PULL_LEVER, PICK_UP_ITEM, USE_TOOL.

Performance Measure

Primary Goal: Reach the finish tile or victory state. Metric of Adaptability: Minimization of the error rate and speed of deduction of the hidden rule in the fewest possible moves/time.

Conclusion: Beyond Imitation Towards True Adaptation

Summary of the ZSDT vs. Turing Test comparison and final reflections

Direct Comparison (ZSDT vs. Turing Test):

- **Goal:** Linguistic deception (Turing) vs. Logical problem solving (ZSDT).
- **Evaluation:** Subjective (easily deceived human judges) vs. Objective (quantitative metrics: time, moves, success rate).
- **Data Dependency:** At risk of memorization/training bias vs. Pure Zero-Shot deduction on unseen rules.
- **Anthropocentrism:** Requires appearing human vs. Requires only demonstrating adaptability and reasoning.

Key Takeaway:

The ability to hold a fluent conversation draws from massive amounts of stored data but does not guarantee any real understanding.

True intelligence does not lie in the ability to imitate humans, but in the ability of an autonomous agent to understand unknown rules, adapt dynamically and deductively to the unexpected, and solve problems in real-time.